

Simple, Fast and Practical Non-blocking and Blocking Concurrent Queue Algorithms

Maged M. Michael
Michael L. Scott

Department of Computer Science
University of Rochester

Presented by: Jun Miao

Outline

- Introduction
- Related work
- 2-lock FIFO queue algorithm
- Non-blocking FIFO queue algorithm
- Experiment
- Conclusion
- Advantages and Disadvantages

Introduction

- FIFO queue: A linked list of items with two interface methods, namely, enqueue and dequeue.
- Problem: How to make enqueues and dequeues work concurrently and correctly on a shared FIFO queue

Some Definitions

- Wait-free
 - All threads or processes can make progress in finite steps
- Non-blocking
 - Some threads or processes are guaranteed to make progress in finite steps
- Obstruction-free
 - One thread or process is guaranteed to make progress in finite steps
- Lock-free
 - No deadlock or Not using locks?


Methods

- Blocking: utilizes locks to prevent conflicts between competing processes and threads
 - Most significant feature: Slow processes and threads may block fast ones
- Non-Blocking: ensures that threads competing for a shared resource do not have their execution indefinitely postponed by mutex or semaphore.

Related Work

- Hwang and Briggs[1984], Sites[1980] and Stones[1990] present lock-free algorithms based on `compare_and_swap`.
- Gottlieb et al.[1983] and Mellor-Crummey[1987] present algorithms that are lock-free but not non-blocking.
- Prakash, Lee and Johnson[1991,1994] present a non-blocking algorithm by taking snapshots of the queue
- Valois[1994,1995] presents a list-based non-blocking algorithm.

The FIFO queue



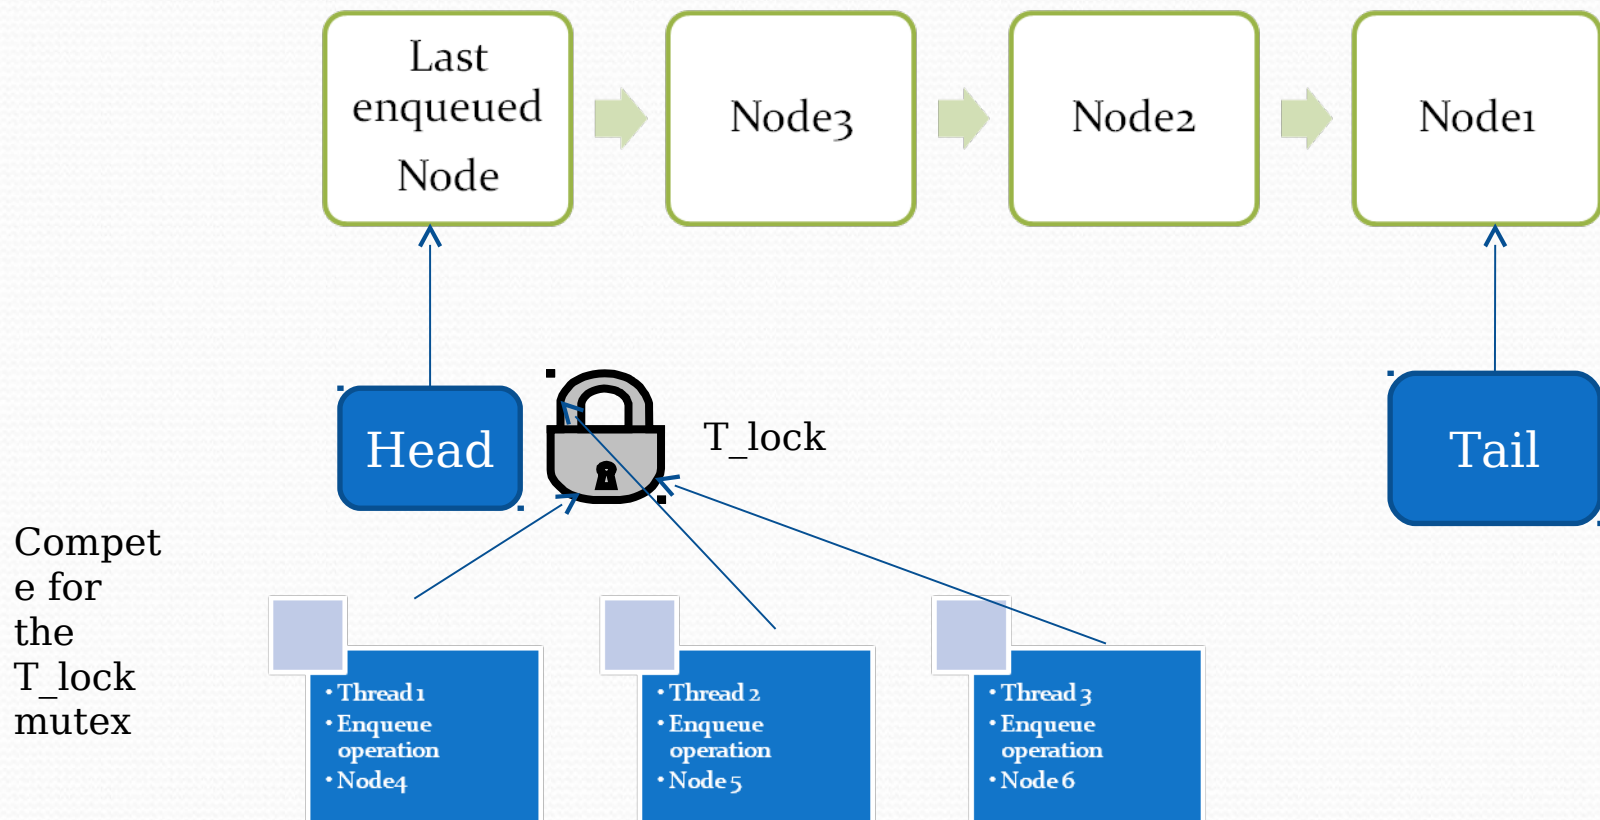
Head: always point to the last dequeued point

Tail: Should point to the last node in the queue

The only inconsistent condition: a new node is enqueued as the last node, but Tail is not updated to point to it.

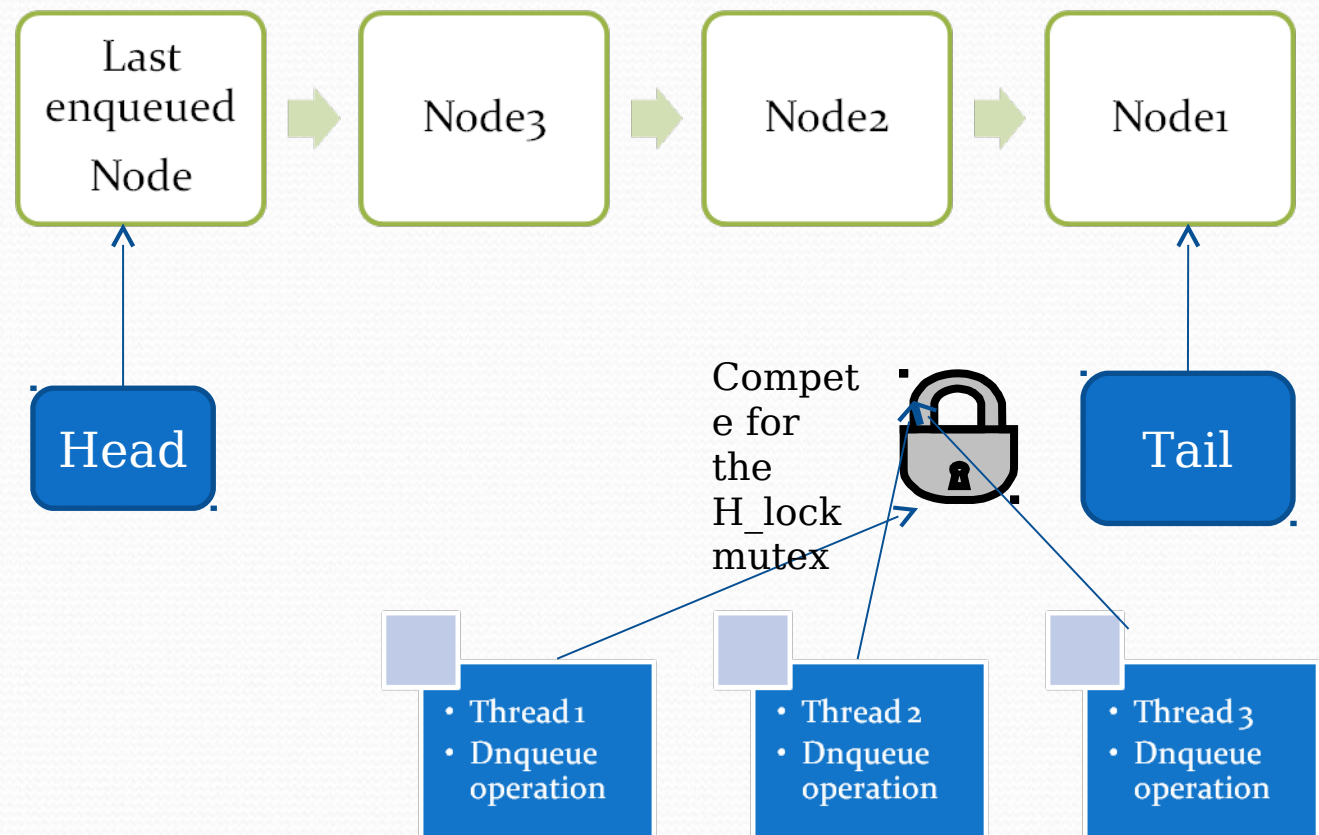
2-Lock concurrent FIFO queue algorithm

- Enqueue operation



2-Lock concurrent FIFO queue algorithm(Cont'd)

- **Dequeue operation**



2-Lock concurrent FIFO queue algorithm (Cont'd)

dequeue(Q: pointer to queue_t, pvalue: pointer to data type): **boolean**

lock(&Q->H_lock)	# Acquire H_lock in order to access Head
node = Q->Head	# Read Head
new_head = node->next	# Read next pointer
if new_head == NULL	# Is queue empty?
unlock(&Q->H_lock)	# Release H_lock before return
return FALSE	# Queue was empty
endif	
*pvalue = new_head->value	# Queue not empty. Read value before release
Q->Head = new_head	# Swing Head to next node
unlock(&Q->H_lock)	# Release H_lock
free(node)	# Free node
return TRUE	# Queue was not empty, dequeue succeeded

Compare And Swap


- A special instruction that atomically compares the contents of a memory location to a given value and, if they are the same, modifies the contents of that memory location to a given new value.

C code of CAS:

```
int compare_and_swap (int *word, int testval, int newval)
{
    int oldval; oldval = *word;
    if (oldval == testval) *word = newval;
    return oldval;
}
```


ABA Problem

- ABA problem is the enemy of non-blocking algorithm using CAS primitive.



ABA Problem(Cont'd)

- Next

A
shared
memory
whose
address
is

Value: A

Thread t2:
Change Value
from A to B, and
then back to A.

- Finally

A
shared
memory
whose
address
is

Value: A

Thread t1:
CAS(&ADDR, A,
newValue)

CAS will succeed because value in the address is still A.
But the environment may change!

Solution

- Associate a counter with each block of memory

`<node, counter>`

If the node changes, `counter++`

- Compare not only the value in the address, but also the counters when using CAS

Non-blocking FIFO queue algorithm

```
structure pointer_t    {ptr: pointer to node_t, count: unsigned integer}  
structure node_t      {value: data type, next: pointer_t}  
structure queue_t     {Head: pointer_t, Tail: pointer_t}
```

```
initialize(Q: pointer to queue_t)  
    node = new_node()           # Allocate a free node  
    node->next.ptr = NULL       # Make it the only node in the linked list  
    Q->Head = Q->Tail = node    # Both Head and Tail point to it
```

```
enqueue(Q: pointer to queue_t, value: data type)  
E1:    node = new_node()        # Allocate a new node from the free list  
E2:    node->value = value       # Copy enqueued value into node  
E3:    node->next.ptr = NULL     # Set next pointer of node to NULL  
E4:    loop                     # Keep trying until Enqueue is done  
E5:        tail = Q->Tail        # Read Tail.ptr and Tail.count together  
E6:        next = tail->next      # Read next ptr and count fields together  
E7:        if tail == Q->Tail    # Are tail and next consistent?  
E8:            if next.ptr == NULL # Was Tail pointing to the last node?  
E9:                if CAS(&tail->next, next, <node, next.count+1>) # Try to link node at the end of the linked list  
E10:                    break     # Enqueue is done. Exit loop  
E11:                endif  
E12:            else             # Tail was not pointing to the last node  
E13:                CAS(&Q->Tail, tail, <next.ptr, tail.count+1>) # Try to swing Tail to the next node  
E14:            endif  
E15:        endif  
E16:    endloop  
E17:    CAS(&Q->Tail, tail, <node, tail.count+1>) # Enqueue is done. Try to swing Tail to the inserted node
```

Non-blocking FIFO queue algorithm(Cont'd)

```
dequeue(Q: pointer to queue.t, pvalue: pointer to data type): boolean
D1:      loop
D2:      head = Q->Head
D3:      tail = Q->Tail
D4:      next = head->next
D5:      if head == Q->Head
D6:          if head.ptr == tail.ptr
D7:              if next.ptr == NULL
D8:                  return FALSE
D9:              endif
D10:             CAS(&Q->Tail, tail, <next.ptr, tail.count+1>)
D11:         else
D12:             # Read value before CAS, otherwise another dequeue might free the next node
D13:             *pvalue = next.ptr->value
D14:             if CAS(&Q->Head, head, <next.ptr, head.count+1>)
D15:                 break
D16:             endif
D17:         endif
D18:     endloop
D19:     free(head.ptr)
D20:     return TRUE
```

Keep trying until Dequeue is done
Read Head
Read Tail
Read Head.ptr->next
Are head, tail, and next consistent?
Is queue empty or Tail falling behind?
Is queue empty?
Queue is empty, couldn't dequeue
Tail is falling behind. Try to advance it
No need to deal with Tail
Try to swing Head to the next node
Dequeue is done. Exit loop
It is safe now to free the old dummy node
Queue was not empty, dequeue succeeded

Non-blocking FIFO queue algorithm(Cont'd)

- For a non-blocking algorithm, every CAS primitive will be executed only after step-by-step checks.
- Try to split the procedure into atomic operations to detect every interleaves

Experiment

- Work on 12-processor Silicon Graphics Challenge multiprocessor
- Associate each processes with each processors
- Totally 100,000 pairs enqueue and dequeue operation.
- $100,000/n$ on each processor. n is the number of processors being used.

Result




Figure 3: Net execution time for one million enqueue/dequeue pairs on a dedicated multiprocessor.

Conclusion

- The concurrent non-blocking queue algorithm that is simple, practical, and fast. It is the choice of any queue-based application on a multiprocessor with universal atomic primitives.
- The 2-lock is recommended to heavily-utilized queues on machines with simple atomic primitives.

Advantages

- Blocking
 - Easy to implement
 - Can be used universally
- Non-blocking
 - High efficiency
 - Never deadlock
 - Ensures the priority of fast threads

Disadvantages

- Blocking

- allows slow threads blocking fast threads
- Deadlock may happen
- Starvation

- Non-blocking

- Hard to design
- depend on what atomic primitives can be provided by systems
- ABA problem is a potential danger for those who use CAS primitive
- Cannot completely avoid starvation



Question time



*Thank you for your
attention!*